

Insurance Pricing and Risk Analytics Project

1. Introduction

1.1 Project Overview

The insurance industry relies heavily on data-driven decision making to accurately assess risk, determine appropriate premiums, and maintain portfolio profitability. The objective of this project was to design and implement a complete insurance analytics pipeline that demonstrates the full lifecycle of insurance data management, from database creation and validation through to actuarial pricing and risk modelling.

The project combines SQL and Python to create a realistic motor insurance portfolio containing customers, vehicles, policies, and claims. The dataset was generated using synthetic data techniques and stored within a relational SQLite database. Following data generation, extensive validation checks were performed to ensure data quality and consistency before exploratory analysis and actuarial modelling were undertaken.

The final stages of the project focused on developing a pricing framework based on claim frequency and claim severity modelling. Additional stress testing and Monte Carlo simulation techniques were then applied to evaluate portfolio risk under adverse scenarios.

1.2 Project Objectives

The key objectives of the project were:

- Design and implement a relational insurance database.
- Generate a realistic insurance portfolio.
- Validate data quality using SQL and Python.
- Identify claim trends and risk factors through exploratory analysis.
- Engineer useful predictive features.
- Develop claim frequency and severity pricing models.
- Estimate pure premiums for individual policyholders.
- Evaluate portfolio performance under stress scenarios.
- Quantify portfolio risk using Monte Carlo simulation.
- Demonstrate advanced SQL analytical techniques.

1.3 Technology Stack

The project utilised the following technologies:

Technology	Purpose
SQLite	Relational database management
SQL	Data validation and analytics
Python	Data cleaning and modelling

Technology Purpose

Pandas	Data manipulation
Statsmodels	GLM actuarial modelling
Matplotlib	Visualisation
Seaborn	Statistical visualisation
Faker	Synthetic data generation

2. Database Design and Data Generation

2.1 Database Structure

The insurance portfolio was modelled using four primary tables:

Customers

The Customers table stores demographic information for policyholders.

Field

CustomerID
Name
Age
Gender
Location
DatePassedDrivingTest

Vehicle

The Vehicle table stores information relating to insured vehicles.

Field

VehicleID
CustomerID
Brand
Age
Mileage
Value

Policy

The Policy table stores policy-level information.

Field

PolicyID

Field

CustomerID

VehicleID

StartDate

EndDate

LengthOfPolicy

CoverageType

Excess

Premium

Claims

The Claims table stores all submitted claims.

Field

ClaimsID

PolicyID

ClaimDate

ClaimType

ClaimStatus

Amount

2.2 Primary and Foreign Keys

To ensure referential integrity, primary and foreign key relationships were established throughout the database.

- CustomerID uniquely identifies each customer.
- VehicleID uniquely identifies each vehicle.
- PolicyID uniquely identifies each policy.
- ClaimsID uniquely identifies each claim.

Foreign key relationships include:

- Vehicle.CustomerID → Customers.CustomerID
- Policy.CustomerID → Customers.CustomerID
- Policy.VehicleID → Vehicle.VehicleID
- Claims.PolicyID → Policy.PolicyID

This structure ensures that all policies are linked to valid customers and vehicles while every claim is linked to a valid policy.

2.3 Data Generation Process

Synthetic data was generated using Python and the Faker library.

The customer population consisted of approximately 10,000 policyholders. Customer ages were generated using a beta distribution to create a realistic concentration of younger and middle-aged drivers while maintaining smaller proportions of older drivers.

Vehicle values were determined using:

- Vehicle brand
- Vehicle age
- Annual mileage
- Depreciation factors

Luxury vehicle brands such as Ferrari, Lamborghini and Porsche were assigned substantially higher base values than standard vehicle brands.

Policy premiums were generated using an expected loss framework based on:

$\text{Expected Loss} = \text{Frequency} \times \text{Severity}$

Additional premium adjustments were applied for:

- Driver age
- Coverage level
- Vehicle value
- Policy excess

Claims were generated probabilistically with approximately 18% of policies experiencing at least one claim.

3. Data Cleaning and Validation

3.1 Validation Framework

A comprehensive validation framework was implemented using SQL queries and Python scripts to ensure the generated dataset remained realistic and internally consistent.

Validation checks focused on:

- Duplicate records
- Missing values
- Referential integrity
- Range validation
- Business rule validation

3.2 Duplicate Record Checks

Duplicate primary keys were investigated across all tables using SQL aggregation queries.

No duplicate primary key values were identified within:

- Customers
- Vehicle
- Policy
- Claims

This confirmed successful implementation of primary key constraints.

3.3 Missing Value Analysis

All critical fields were assessed for missing values.

The following attributes were validated:

- Customer demographics
- Vehicle characteristics
- Policy information
- Claims information

No missing values were detected within mandatory fields.

3.4 Customer Validation

Customer ages were assessed to ensure values remained within realistic boundaries.

Validation confirmed:

- Minimum age exceeded legal driving age requirements.
- No customers exceeded the maximum expected age threshold.
- Driver test dates remained consistent with customer ages.

3.5 Vehicle Validation

Vehicle values and mileage records were validated.

Checks confirmed:

- No vehicles possessed negative mileage.
- No vehicles possessed zero or negative values.
- Vehicle values remained consistent with depreciation assumptions.

3.6 Policy Validation

Policy start and end dates were validated.

Checks confirmed:

- Policy end dates always occurred after policy start dates.
- Policy durations were correctly calculated.
- Premiums remained within expected commercial ranges.

3.7 Claims Validation

Claims data was validated using multiple business rules.

Checks confirmed:

- No negative claim amounts.
- No denied claims received positive payments.
- Claim payments never exceeded vehicle values.
- Approved claims followed coverage restrictions.

These checks provided confidence that the claims portfolio accurately reflected realistic insurance behaviour.

4. Exploratory Data Analysis

4.1 Introduction

Following successful validation of the dataset, exploratory data analysis (EDA) was performed to identify patterns within the insurance portfolio. The objective of this stage was to understand the distribution of claims, identify high-risk customer segments, and investigate how claim behaviour varied across different demographic and policy characteristics.


The analysis focused on:

- Claim type distributions
- Customer age profiles
- Vehicle brand performance
- Claims trends over time
- Portfolio profitability indicators

4.2 Claim Type Distribution

The first stage of the analysis examined how total claim costs were distributed across different claim categories.

Figure 1: Distribution of Total Claim Amount by Claim Type



The results show that theft-related claims represented the largest proportion of total claim costs, accounting for 28.81% of all claims (£4.35 million). Fire claims were the second largest contributor, representing 27.69% (£4.18 million), followed closely by major accidents at 24.03% (£3.63 million).

While weather accounted for the most number of claims, it only amounted 3.91% of the total claim amount.

These findings indicate that although theft and fire claims occur less frequently than minor accidents, their significantly higher severity causes them to dominate overall claim expenditure. This observation supports the need for severity modelling within the pricing framework, as claim frequency alone would underestimate portfolio risk.


Key Findings

- Theft generated the largest share of losses.
- Fire claims represented a substantial source of severity risk.
- Major accidents contributed almost one-quarter of all claim costs.
- Lower-severity claim types had limited impact on overall portfolio losses.

4.3 Claim Amount by Customer Age

The relationship between age and total claim cost was investigated to determine whether particular age groups represented elevated levels of risk.

Figure 2: Total Claim Amount by Customer Age



The analysis revealed a clear concentration of claim costs among younger drivers. Total claims peaked at age 26, reaching approximately £825,000. Following this peak, claim amounts gradually declined as driver age increased.

The distribution was noticeably right-skewed, with the majority of claims concentrated between ages 25 and 40. Drivers aged over 60 generated the lowest total claim costs across the portfolio.

This trend is consistent with established insurance theory, where younger drivers generally possess less driving experience and therefore exhibit higher accident frequencies and claim costs.


Key Findings

- Age 26 generated the highest total claim amount.
- Most claims were concentrated between ages 25 and 40.
- Older drivers exhibited significantly lower claim costs.
- Age appears to be an important rating factor for pricing purposes.

4.4 Claim Amount by Vehicle Brand

Vehicle brands were analysed to determine whether certain manufacturers generated disproportionately high claim costs.

Figure 3: Total Claim Amount by Vehicle Brand



The results demonstrate substantial variation between vehicle brands.

Porsche generated the highest total claim costs, followed by Ferrari and Lamborghini. These luxury vehicle manufacturers produced claim totals ranging between approximately £2 million and £3 million.

In contrast, Volkswagen, Ford and Peugeot generated the lowest claim costs within the portfolio.

Most vehicle brands generated total claims between £500,000 and £1 million, suggesting relatively stable loss experience among standard vehicle categories. The substantially higher losses associated with luxury vehicles likely reflect their greater repair costs, higher replacement values and increased claim severity.


Key Findings

- Porsche represented the highest-risk vehicle brand.
- Ferrari and Lamborghini also generated exceptionally high losses.
- Standard vehicle brands produced relatively stable claim costs.
- Luxury vehicles appear significantly more expensive to insure.

4.5 Claims Over Time

Claim frequency was analysed over time to identify any emerging trends or unusual patterns.

Figure 4: Number of Claims by Year



Annual claim frequency remained relatively stable throughout most of the observation period, averaging approximately 40 claims per year.

The highest claim frequency occurred in 2004, when 73 claims were recorded. The lowest frequency occurred in 2003, with 27 claims.

A noticeable decline occurs after 2026. This reduction is attributable to the synthetic nature of the dataset and the policy generation process rather than an actual reduction in risk.

Overall, the relatively stable claim frequency suggests that the portfolio does not exhibit substantial temporal volatility.


Key Findings

- Claims remained broadly stable over time.
- Peak frequency occurred in 2004.
- Lowest frequency occurred in 2003.
- No sustained upward or downward trend was observed.

4.6 Portfolio Profitability Analysis

The relationship between premiums collected, claims paid and resulting profit was examined to assess portfolio performance.

Figure 5: Net Profit, Earned Premium and Claims Paid



The analysis indicates that the portfolio remained profitable for the majority of the observation period.

Net profit reached its highest level in 2009, generating approximately £350,000. The largest loss occurred in 2011, where portfolio profitability fell to approximately negative £70,000.

Additional periods of negative profitability were observed during parts of the 1970s. However, these periods were relatively infrequent and limited in magnitude.

Overall, earned premium generally exceeded claims paid, demonstrating that the pricing framework generated a sustainable level of profitability.

Key Findings


- Portfolio profitability was positive in most years.
- Largest profit occurred in 2009.
- Largest loss occurred in 2011.
- Premium income generally exceeded claims expenditure.

4.7 Loss Ratio Analysis

Loss ratio represents one of the most important performance indicators within insurance pricing.

Loss Ratio = Claims Paid / Premium Earned

Figure 6: Loss Ratio by Year



Loss ratios ranged between approximately 0.10 and 1.18 throughout the analysis period.

The highest loss ratio values occurred during the same periods where profitability was lowest, while the lowest loss ratios coincided with the most profitable years.

A loss ratio above 1 indicates that claims exceeded premium income, resulting in underwriting losses. Conversely, loss ratios below 1 indicate profitable underwriting performance.

The majority of observations remained below 1, supporting the conclusion that the portfolio was generally profitable.

Key Findings

- Loss ratios ranged between 0.10 and 1.15.
- Most years remained below the critical threshold of 1.
- Higher loss ratios corresponded with reduced profitability.
- Portfolio pricing remained broadly adequate.

4.8 Summary of Exploratory Analysis

The exploratory analysis identified several important risk characteristics within the portfolio.

The strongest drivers of claim cost were:

1. Vehicle brand.
2. Customer age.
3. Claim type.
4. Coverage level.

Luxury vehicles generated significantly higher losses than standard vehicles, while younger drivers consistently produced larger claim totals than older drivers. Theft, fire and major accident claims dominated portfolio losses due to their high severity.

These findings provide a strong foundation for the actuarial pricing models developed in the next stage of the project.

5. Feature Engineering and Pricing Model Development

5.1 Introduction

Following exploratory analysis, additional predictive variables were created to improve the quality of the actuarial pricing models. Feature engineering is a critical stage within insurance analytics because raw transactional data often does not directly capture underlying risk characteristics.

The objective of this stage was to derive meaningful indicators of policyholder behaviour and claims experience which could subsequently be incorporated into pricing models.

The engineered features were created using SQL views and stored within the database to allow efficient retrieval and future reuse.

5.2 Feature Engineering

Several new variables were developed from the underlying claims and policy data.

Customer Claim Frequency

Customer claim frequency measures the total number of claims submitted by each customer.

SQL View

```
CREATE VIEW customer_claim_frequency AS  
  
SELECT  
  
    p.customerid,  
  
    COUNT(c.claimsid) AS claim_frequency  
  
FROM policy p  
  
INNER JOIN claims c  
  
    ON p.policyid = c.policyid  
  
GROUP BY p.customerid;
```

Claim frequency is widely used within actuarial pricing because historical claims behaviour often provides insight into future risk.

Customer Claim Summary

A customer-level summary was created to consolidate overall claims performance.

The summary included:

- Total number of claims
- Total claim cost
- Average claim cost

This view provides a concise measure of individual policyholder risk.

SQL View

```
CREATE VIEW customer_claim_summary AS
```

```
SELECT
```

```
    p.customerid,
```

```
    COUNT(c.claimsid) AS number_of_claims,
```

```
    SUM(c.amount) AS total_claim_amount,
```

```
    AVG(c.amount) AS avg_claim_amount
```

```
FROM policy p
```

```
INNER JOIN claims c
```

```
    ON p.policyid = c.policyid
```

```
GROUP BY p.customerid;
```

Policy Claim Summary

A policy-level view was created to measure claims experience for each individual policy.

SQL View

```
CREATE VIEW policy_claim_summary AS
```

```
SELECT
```

```
    policyid,
```

```
    COUNT(claimsid) AS number_of_claims,
```

```
    SUM(amount) AS total_claim_amount,
```

```
    AVG(amount) AS avg_claim_amount
```

```
FROM claims
```

```
GROUP BY policyid;
```

This view provides a convenient summary for underwriting and portfolio monitoring purposes.

Average Claim Severity

Average claim severity was calculated across the portfolio.

SQL View

```
CREATE VIEW average_claim_severity AS
```

```
SELECT
```

```
    AVG(amount) AS avg_claim_severity
```

```
FROM claims;
```

Claim severity represents the average financial impact of a claim and forms one of the two core components of actuarial pricing.

5.3 Pricing Framework

The pricing methodology adopted within this project follows a standard actuarial pure premium approach.

Expected claim costs are estimated by combining:

- Claim frequency
- Claim severity

The expected loss is therefore calculated as:

Expected Loss = Claim Frequency x Claim Severity

This approach separates the probability of a claim occurring from the expected cost of that claim, allowing each component to be modelled independently.

6. Claim Frequency Modelling

6.1 Poisson Regression Model

Claim frequency represents count data and therefore lends itself naturally to Poisson regression techniques.

A Generalised Linear Model (GLM) with a Poisson distribution was fitted using:

- Age
- Gender
- Coverage type

as explanatory variables.

The model estimates the expected number of claims for each policyholder.

Model Inputs

Variable	Type
Age	Numeric
Gender	Categorical
Coverage Type	Categorical

The target variable was:

Number of Claims

6.2 Frequency Model Results

The Poisson model successfully estimated claim frequencies across the portfolio.

A dispersion analysis was also performed to assess whether the Poisson assumption was appropriate.

The variance of claim counts was found to be reasonably consistent with the mean, suggesting that the Poisson framework provided an acceptable representation of claim frequency behaviour.

Figure 7: Frequency Model Summary Output

```

=====
Generalized Linear Model Regression Results
=====
Dep. Variable:      claim_count    No. Observations:      13001
Model:              GLM           Df Residuals:           12996
Model Family:       Poisson       Df Model:               4
Link Function:      Log           Scale:                 1.0000
Method:             IRLS          Log-Likelihood:        -8399.4
Date:               Thu, 11 Jun 2026 Deviance:               11674.
Time:               16:04:43       Pearson chi2:          2.08e+04
No. Iterations:     6             Pseudo R-squ. (CS):    0.0005899
Covariance Type:    nonrobust
=====

              coef      std err      z      P>|z|      [0.025      0.975]
-----
const          -2.0756      0.074    -28.106      0.000     -2.220     -1.931
age            -0.0028      0.002     -1.522      0.128     -0.006      0.001
gender_M        0.0726      0.036      2.023      0.043      0.002      0.143
coverage_Third party  0.0455      0.046      0.992      0.321     -0.044      0.135
coverage_Third party, fire and theft  0.0017      0.044      0.040      0.968     -0.084      0.087
=====

===== DISPERSION CHECK =====

Mean Claim Count: 0.23905853395892623
Variance Claim Count: 0.304385082804236
Poisson variance assumption appears reasonable.

===== RISK MULTIPLIERS =====

const          0.125481
age            0.997210
gender_M       1.075326
coverage_Third party  1.046501
coverage_Third party, fire and theft  1.001730

```

The model identified meaningful relationships between customer characteristics and expected claim frequency.

Younger drivers generally exhibited higher predicted claim frequencies, while coverage type also influenced expected claim counts.

6.3 Frequency Predictions

Predicted claim frequencies were generated for every customer within the portfolio.

These predictions form the first component of the pure premium calculation.

7. Claim Severity Modelling

7.1 Gamma Regression Model

Claim severity was modelled using a Gamma Generalised Linear Model with a logarithmic link function.

The Gamma distribution is widely used in insurance because claim amounts are:

- Positive
- Continuous
- Right-skewed

The model used:

- Age
- Gender
- Coverage Type
- Number of Claims

as explanatory variables.

7.2 Severity Model Results

The severity model estimated expected claim costs for individual customers.

Figure 8: Gamma Severity Model Summary

```

Generalized Linear Model Regression Results
=====
Dep. Variable:      avg_claim_amount    No. Observations:      1503
Model:              GLM                Df Residuals:          1497
Model Family:       Gamma              Df Model:              5
Link Function:      Log                Scale:                2.8619
Method:             IRLS              Log-Likelihood:        -15139.
Date:               Thu, 11 Jun 2026    Deviance:              2655.2
Time:               16:22:08           Pearson chi2:          4.28e+03
No. Iterations:     12                 Pseudo R-squ. (CS):    0.01959
Covariance Type:    nonrobust
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	9.1149	0.228	40.018	0.000	8.668	9.561
age	-0.0006	0.005	-0.124	0.901	-0.010	0.009
number_of_claims	-0.1285	0.103	-1.253	0.210	-0.330	0.072
gender_M	-0.1379	0.087	-1.579	0.114	-0.309	0.033
coverage_type_Third party	-0.1362	0.126	-1.080	0.280	-0.384	0.111
coverage_type_Third party, fire and theft	0.4020	0.098	4.098	0.000	0.210	0.594

```

=====
===== SEVERITY RISK MULTIPLIERS =====
const                9089.922360
age                  0.999426
number_of_claims     0.879384
gender_M             0.871160
coverage_type_Third party    0.872648
coverage_type_Third party, fire and theft 1.494741

```

The model demonstrated expected insurance behaviour, with claim severity increasing for certain customer and policy characteristics.


Because very large claims occur relatively infrequently, some prediction error remained within the highest severity observations.

This behaviour is common within insurance modelling and reflects the difficulty of predicting rare catastrophic losses.

7.3 Actual vs Predicted Severity

The predictive performance of the severity model was assessed by comparing actual claim amounts against predicted values.

Figure 9: Actual vs Predicted Claim Severity



The results indicate that the model captures overall severity patterns reasonably well.

Predictions generally increase alongside actual claim amounts, demonstrating that the model successfully identifies broad severity trends.

However, the model tends to underestimate a small number of extremely large claims. This is a common characteristic of insurance pricing models because catastrophic claims occur infrequently and contain limited historical information.

Key Findings

- Strong relationship between actual and predicted values.
- Good performance for typical claims.
- Slight underestimation of extreme losses.
- Suitable for portfolio-level pricing applications.

8. Pure Premium Modelling

8.1 Pure Premium Calculation

The final pricing metric was calculated by combining frequency and severity predictions.


Pure Premium = Predicted Frequency × Predicted Severity

This represents the expected annual claim cost for each policyholder.

8.2 Pure Premium Distribution

The distribution of predicted pure premiums was examined to understand overall portfolio risk.

Figure 10: Pure Premium Distribution



The distribution displays a clear concentration of customers within lower and medium-risk categories, with a smaller proportion of customers exhibiting substantially higher expected losses.

This pattern is typical of insurance portfolios where a relatively small number of policyholders account for a disproportionately large share of expected claims costs.


8.3 Risk Band Classification

Customers were segmented into three risk categories:

- Low Risk
- Medium Risk
- High Risk

using quantile-based thresholds.

Figure 11: Pure Premium by Risk Band



The resulting segmentation clearly differentiates customers according to expected loss.

Low-risk customers possess significantly lower expected costs, while high-risk customers generate considerably larger expected losses.

This segmentation could be used directly within a commercial pricing framework to support underwriting decisions and premium differentiation.

Key Findings

- Clear separation between risk groups.
- High-risk customers exhibit substantially larger expected losses.
- Pure premium provides an effective measure of customer risk.

- Risk segmentation supports targeted pricing strategies.

8.4 Pricing Model Summary

The pricing framework successfully combines claim frequency and claim severity modelling to estimate expected future losses.

The resulting pure premium metric provides a practical foundation for insurance pricing and demonstrates how actuarial modelling techniques can be applied to real-world insurance portfolios.

The next stage of the project investigates how this pricing framework performs under adverse scenarios and extreme events.

9. Advanced SQL Analytics

9.1 Introduction

The purpose of the advanced SQL analytics stage was to move beyond simple aggregation and explore more sophisticated methods of identifying risk patterns within the insurance portfolio.

Window functions were used to rank customers, age groups, policy characteristics and coverage types according to multiple risk measures, whilst complex joins allowed claims experience to be analysed alongside policyholder and policy information.

These techniques replicate many of the analytical processes used within actuarial and insurance reporting environments.

9.2 Customer Risk Ranking

Customers were ranked using three separate measures:

- Total number of claims
- Total claim cost
- Average claim severity

The objective was to identify policyholders contributing the greatest level of portfolio risk.

Figure 12: Top Customers by Total Claims

CustomerID	Name	Age	Location	Gender	claim_frequency	risk_rank
1642	Justin Johnson	23	Lake Kimberly	M	4	1
1759	Linda Vance	33	East Thomas	F	4	1
1833	Anthony Carr	30	Port Emilyfort	M	4	1
1844	Matthew Navarro	34	East Brandon	M	4	1
3194	Evelyn Johnson	37	Bridgetstad	F	4	1
3528	Angel Nash	27	Thomaston	M	4	1

Customers with multiple claims consistently appeared at the top of the ranking. This suggests that previous claims behaviour may be a useful indicator of future claims activity.

Figure 13: Top Customers by Total Claim Cost

CustomerID	Name	Age	Location	Gender	claim_total	risk_rank
2545	Elizabeth Ramos	34	Jenkinsview	F	170888.26	1
8102	Joseph Barrett	38	Ericksonmouth	M	150355.0	2
3501	Kristina Harris	37	West Monica	F	149573.19	3
7365	Derek Lowe	26	Williamsville	M	146392.04	4
3526	John Wright	41	Port Priscilla	M	138864.93	5
2555	Sarah Gonzalez	30	Jennifertown	F	128882.9	6

A relatively small number of customers generated substantially higher losses than the remainder of the portfolio.

This finding supports a common insurance principle that losses are often concentrated amongst a small subset of policyholders.

Figure 14: Top Customers by Average Claim Severity

CustomerID	Name	Age	Location	Gender	number_of_claims	avg_claim_amount
8102	Joseph Barrett	38	Ericksonmouth	M	1	150355.0
3501	Kristina Harris	37	West Monica	F	1	149573.19
5360	Tamara Johnson	32	Schultzside	F	1	107441.19
7128	Jennifer Long	20	Lake Johnmouth	F	1	104417.19
2836	Dawn Lopez	26	New Elizabethburgh	F	1	100956.0

Customers exhibiting high average claim severity represent a significant source of underwriting risk, even if claim frequency remains relatively low.

This demonstrates the importance of modelling both frequency and severity separately within actuarial pricing.

9.3 Age-Based Risk Analysis

Age is one of the most widely used rating factors within motor insurance because driver behaviour often changes throughout the lifecycle.

Age groups were ranked using:

- Claim frequency
- Total claims generated
- Loss ratio

Figure 15: Claim Amount by Age

Age	claim_total	risk_rank
26	836284.53	1
28	774139.95	2
33	728391.42	3
25	671681.32	4
36	654823.32	5
37	620417.68	6

The analysis shows a pronounced concentration of claim costs amongst younger drivers.

Claim amounts peak at approximately age 26, where total claims exceed £825,000.

Following this peak, claim costs gradually decline as age increases.

The distribution is positively skewed, with most claims occurring between ages 25 and 40, whilst customers aged over 60 generate comparatively low levels of claims activity.

This result is consistent with established insurance experience, where younger drivers generally exhibit higher risk due to reduced driving experience and more frequent accident involvement.

9.4 Coverage Type Risk Analysis

Coverage types were analysed to determine whether certain products generated greater claims experience than others.

The analysis compared:

- Claim frequency
- Claim severity
- Loss ratio

Figure 16: Coverage Type Risk Ranking

CoverageType	average_claim_frequency	risk_rank
Third party, fire and theft	1201.0	1
Third party	963.0	2
Fully Comprehensive	944.0	3

Coverage products offering broader protection generally generated higher claim frequencies and greater average claim severities.

This outcome is expected because comprehensive policies cover a wider range of loss events than third-party products.

The findings demonstrate how product design directly influences portfolio risk exposure.

9.5 Relative Risk Analysis

Relative risk analysis was performed to identify groups performing above or below the portfolio average.

The analysis compared:

- Customers versus average customer loss
- Age groups versus portfolio average loss
- Coverage types versus baseline loss

Figure 17: Relative Risk Classification

customerid	total_loss	avg_loss	risk_flag
2545	170888.26	6764.97242492156	Above Average Risk
8102	150355.0	6764.97242492156	Above Average Risk
3501	149573.19	6764.97242492156	Above Average Risk
7365	146392.04	6764.97242492156	Above Average Risk
3526	138864.93	6764.97242492156	Above Average Risk
2555	128882.9	6764.97242492156	Above Average Risk

The analysis successfully identified customer segments operating above the portfolio average risk level.

These segments could be targeted for:

- Premium adjustments
- Enhanced underwriting review
- Additional risk management controls

9.6 Claim Intensity Analysis

Policy duration was investigated to determine whether longer policy periods generated greater claims activity.

Figure 18: Claims by Policy Duration

LengthOfPolicy	claim_frequency	risk_rank
2	1052	1
1	1035	2
3	1021	3

LengthOfPolicy	claim_amount	risk_rank
2	5148532.15	1
3	5143519.2	2
1	4800602.13	3

The results are very equal and there does not seem to be any correlation between policy length and the claim frequency/severity.

9.7 Rolling Average Analysis

A rolling average of claims was calculated using SQL window functions.

Figure 19: Rolling Average Claims

claim_year	Amount	rolling_avg_claim
1963	5952.62	5952.62
1964	0.0	2976.31
1965	3772.06	3241.56
1966	0.0	1257.3533333333333
1967	2642.97	2138.3433333333333
1968	0.0	880.99

The rolling average smooths short-term fluctuations and highlights longer-term trends within claims behaviour.

This technique is commonly used in insurance reporting to distinguish underlying trends from random year-to-year variation.

10. Stress Testing and Scenario Analysis

10.1 Introduction

Insurance companies must understand how changes in claims behaviour impact profitability.

To assess portfolio resilience, several stress-testing scenarios were applied to the pricing model.

The scenarios examined:

- Increased claim frequency
- Increased claim severity
- Catastrophic events


10.2 Frequency Stress Testing

The first scenario increased expected claim frequencies by:

- 10%
- 20%
- 50%

Expected losses were recalculated under each scenario.

Figure 20: Frequency Stress Testing Results



As expected, portfolio losses increased proportionally with claim frequency.

A 10% increase in claim frequency produced an approximately 10% increase in expected losses.


The results demonstrate the sensitivity of insurer profitability to changes in claim occurrence rates.

Even relatively small increases in frequency can have a significant financial impact when applied across an entire portfolio.

10.3 Severity Stress Testing

A separate scenario increased claim severity by 25%.

Figure 21: Severity Stress Testing



The distribution maintained a similar shape to the baseline model but shifted noticeably upward.

The peak of the distribution increased from approximately £650 to around £800, reflecting the direct effect of larger average claim costs.

The results demonstrate that claim severity inflation can significantly increase expected losses, even when claim frequency remains unchanged.

11. Catastrophe Modelling

11.1 Introduction

Catastrophic events represent low-frequency, high-severity occurrences capable of generating substantial losses across an insurance portfolio.


Examples include:

- Severe storms
- Flooding
- Major fire events
- Widespread theft incidents

To simulate this type of event, 5% of customers were randomly selected and their predicted claim severity was increased by 400%.

11.2 Catastrophe Scenario Results

Figure 22: Catastrophe Scenario Distribution



The catastrophe scenario produced a significant upward shift in portfolio losses.

Portfolio Impact

Metric	Value
Baseline Loss	£15,502,664
Catastrophe Loss	£18,672,005
Increase	£3,169,341

The catastrophe event increased expected portfolio losses by over £3.1 million.

Although only a small proportion of customers were affected, the large increase in severity generated a substantial financial impact.

This highlights the importance of catastrophe modelling, capital management and reinsurance arrangements within insurance operations.

12. Monte Carlo Simulation

12.1 Introduction

Whilst stress testing evaluates specific scenarios, Monte Carlo simulation allows the full distribution of potential outcomes to be examined.

The simulation combines:


- Poisson claim frequency
- Gamma claim severity

to generate thousands of possible future portfolio loss outcomes.

A total of 10,000 simulations were performed.

12.2 Simulation Results

Figure 23: Monte Carlo Portfolio Loss Distribution



The distribution exhibits the expected bell-shaped pattern, centred around the portfolio's expected loss level.

Key Results

Measure	Value
---------	-------

Mean Loss	£15,510,454
-----------	-------------

95% VaR	£16,242,516
---------	-------------

99% VaR	£16,548,647
---------	-------------

12.3 Value at Risk Analysis

Value at Risk (VaR) estimates the loss level expected to be exceeded only within a specified percentage of outcomes.

The results indicate:

- 95% of simulated outcomes fall below approximately £16.24 million.
- 99% of simulated outcomes fall below approximately £16.55 million.

This suggests that extreme portfolio losses remain reasonably contained relative to the expected loss level.

The simulation therefore indicates that the portfolio exhibits a moderate level of risk under the model assumptions.

12.4 Importance of Monte Carlo Modelling

Monte Carlo simulation provides several advantages:

- Captures uncertainty in future claims experience.
- Quantifies tail risk.
- Supports capital modelling.
- Provides insight into extreme loss scenarios.

These techniques are widely used within actuarial reserving, pricing and enterprise risk management.

13. Conclusions and Recommendations

13.1 Project Summary

The aim of this project was to design and develop a complete insurance analytics and pricing framework using SQL and Python. The project replicated many of the processes undertaken within actuarial and insurance data teams, including database design, data validation, exploratory analysis, feature engineering, predictive modelling and risk assessment.

A relational database containing customer, vehicle, policy and claims information was created and populated with over 10,000 customers and associated policy records. SQL was used extensively throughout the project to manage, validate and analyse the data, while Python was used for statistical modelling, simulation and visualisation.

The project successfully demonstrated how raw insurance data can be transformed into actionable insights capable of supporting pricing and risk management decisions.

13.2 Key Findings

Claims Analysis

Analysis of claim types revealed that theft and fire claims generated the largest proportion of losses within the portfolio.

- Theft accounted for approximately 28.8% of all claim costs.
- Fire accounted for approximately 27.7%.
- Major accidents represented approximately 24.0%.

Together, these three claim categories generated more than 80% of all claim costs, indicating that a relatively small number of severe claim types dominate portfolio losses.

Driver Age and Risk

The analysis identified a clear relationship between age and claims experience.

Claim costs peaked at age 26, where total claims exceeded £825,000. Claim amounts remained elevated throughout the 25–40 age range before gradually declining with age. Drivers aged over 60 generated substantially lower claim costs.

These findings support established insurance industry evidence that younger drivers typically represent a higher-risk customer segment.

Vehicle Characteristics

Luxury vehicle brands such as Porsche, Ferrari and Lamborghini generated the highest total claim costs across the portfolio.

In contrast, brands such as Volkswagen, Ford and Peugeot generated significantly lower claims totals.

This suggests that vehicle value and repair costs play an important role in determining overall claims severity.

Pricing Model Performance

The pricing framework successfully combined:

- Poisson frequency modelling
- Gamma severity modelling

to estimate expected future losses.

The resulting pure premium model produced clear separation between low-risk, medium-risk and high-risk customers.

The actual versus predicted severity analysis showed that the model captured overall claims behaviour effectively, although it tended to underestimate a small number of extremely severe claims. This limitation is common within insurance pricing models because very large losses occur infrequently and are therefore difficult to predict accurately.

13.3 Stress Testing Results

Stress testing demonstrated the sensitivity of portfolio performance to changes in claims behaviour.

Increasing claim frequency produced a near proportional increase in expected losses across the portfolio.

Similarly, increasing claim severity by 25% shifted the pure premium distribution upwards and increased expected claim costs across all customer segments.

These results highlight the importance of maintaining adequate pricing margins to absorb unexpected changes in claims experience.

13.4 Catastrophe Modelling Results

The catastrophe scenario simulated a severe event affecting 5% of customers and increasing claim severity by 400%.

The results showed:

Metric	Value
Baseline Portfolio Loss	£15.50 million
Catastrophe Portfolio Loss	£18.67 million
Increase	£3.17 million

Although only a small proportion of policyholders were affected, the financial impact on the portfolio was substantial.

This demonstrates the importance of catastrophe modelling within insurance pricing and highlights why insurers rely on capital reserves and reinsurance arrangements to protect against extreme events.

13.5 Monte Carlo Simulation Results

Monte Carlo simulation provided a probabilistic assessment of future portfolio performance.

The simulation produced:

Measure	Value
Mean Portfolio Loss	£15.51 million
95% VaR	£16.24 million
99% VaR	£16.55 million

The relatively narrow difference between the mean loss and extreme percentile losses suggests that the portfolio remains reasonably stable under the assumptions used within the model.

The simulation demonstrates how actuarial techniques can be used to quantify uncertainty and support risk management decisions.

13.6 Limitations

Several limitations should be acknowledged.

Firstly, the dataset was synthetically generated rather than obtained from a real insurer. While realistic distributions were incorporated, actual insurance portfolios may exhibit more complex behavioural patterns.

Secondly, the pricing model used a relatively limited set of explanatory variables. Commercial pricing models would typically include additional rating factors such as:

- Driving history
- Vehicle usage
- Geographic risk scores
- Previous convictions
- No-claims discounts

Thirdly, the severity model showed some underestimation of extreme losses. More advanced modelling techniques, such as Gradient Boosting Machines or Extreme Value Theory approaches, may improve predictive performance for large claims.

Finally, the catastrophe scenario was based on simplified assumptions and therefore should be viewed as illustrative rather than representative of real-world catastrophe exposure.

13.7 Recommendations

Based on the findings of this project, several recommendations can be made.

Recommendation 1: Enhance Data Collection

Additional customer and vehicle characteristics should be collected to improve pricing accuracy and strengthen predictive performance.

Recommendation 2: Develop More Advanced Pricing Models

Future work could compare traditional actuarial GLMs with machine learning approaches such as:

- Random Forests
- XGBoost
- LightGBM

to evaluate whether predictive accuracy can be improved.

Recommendation 3: Expand Risk Monitoring

The advanced SQL analytics developed within this project could be incorporated into a reporting dashboard to provide continuous monitoring of:

- High-risk customers
- Loss ratios
- Emerging claims trends
- Portfolio profitability

Recommendation 4: Introduce Reinsurance Modelling

Future analysis could incorporate reinsurance structures to evaluate how catastrophic losses would be shared between insurers and reinsurers.

13.8 Final Conclusion

This project successfully developed an end-to-end insurance analytics and pricing framework using SQL and Python. Through database design, data validation, exploratory analysis, feature engineering, actuarial pricing models and risk simulations, the project demonstrated how insurance data can be transformed into meaningful business insights.

The results identified key risk drivers across claim types, customer demographics and vehicle characteristics while demonstrating the practical application of actuarial modelling techniques such as Poisson frequency models, Gamma severity models and Monte Carlo simulation.

Overall, the project provides a strong demonstration of both data analytics and actuarial modelling skills and reflects many of the core processes used within modern insurance pricing and risk management environments.